

Table of contents

- [Table of contents](#)
- [Linux basic operations day 2](#)
 - [Local DNS](#)
- [Linux Shell Scripts](#)
 - [Creating our first shell script](#)
 - [Accepting input from the user](#)
- [Scaling the storage capacity of the instance](#)

Linux basic operations day 2

- Agenda for today would be understanding and performing some of the basic linux operations that will help our understanding

Local DNS

- DNS is domain name service which helps us define a name for an ip address or url
- linux has /etc/hosts which is a DNS file and it is the source of truth for the machine
- Any request originating from the machine will first go to this DNS for resolution , if it doesnt find it then it will go to it's ISP's DNS.
- You can see the external dns in [/etc/resolv.conf](#) in "nameserver property "
- For testing /etc/hosts , we will add few entries in it to test lie below

```
127.0.0.1    localhost localhost.localdomain localhost4 localhost4.localdomain4
::1         localhost6 localhost6.localdomain6
10.0.0.1    www.google.com
10.0.1.25   db1
~
~
```

- Once above entries are added to /etc/hosts , try pinging the server names

```
[root@ip-172-31-87-91 ~]# vi /etc/hosts
[root@ip-172-31-87-91 ~]# ping www.google.com
PING www.google.com (10.0.0.1) 56(84) bytes of data.
^C
--- www.google.com ping statistics ---
3 packets transmitted, 0 received, 100% packet loss, time 2042ms

[root@ip-172-31-87-91 ~]# ping db1
PING db1 (10.0.1.25) 56(84) bytes of data.
^C
--- db1 ping statistics ---
3 packets transmitted, 0 received, 100% packet loss, time 2041ms

[root@ip-172-31-87-91 ~]#
```

- As seen above , even though previously we were able to ping google.com, since we added the entry in local dns , it takes first precedence
- db1 name is also resolved with the ip that we have given

Linux Shell Scripts

Creating our first shell script

- Shell script is nothing but a set of linux shell commands written in an order to execute a task
- We will first use our known commands to create our first shell script
- Let us create

```
sudo su -  
vi install.sh
```

- Once the file opens paste below scripts inside the file

```
#!/bin/bash  
yum install httpd -y  
echo "this is my first script " >> /var/www/html/index.html  
service httpd start
```

```
#once we have pasted the content , run below commands  
chmod 777 install.sh  
#above command will give the fule executable permissions  
./install.sh  
#above command will execute the script , once executed observe the output  
# You can even store the output of the execution in a file using ./install.sh >> output.txt
```

-

Accepting input from the user

- We may feel the need from time and time again to get inputs from the user
- We can fetch the arguments which are passed while executing the script
- for ex ``./test.sh first second`and here`first`and`second`is an argument which is passed to`test.sh`
- Similarly one can pass multiple arguments

```
#!/bin/bash  
echo "the name of script is $0"  
echo "first parameter is $1"  
echo "second parameter is $2"
```

- While executing above script we need to pass 1 argument
- This argument will be printed in the output
- These arguments are also called as positional arguments . if there are 2 arguments which are passed we can use them with \$2 , \$3 etc

```
#!/bin/bash
echo "first number you have entered is " $1
echo "second number you have entered is " $2
sum=$(( $1 + $2 ))
echo "addition of numbers is $sum"
```

- Above script expects 2 positional arguments to be passed while executing a script
- We can also prompt the user to give output that we need
- this output we can read in the script and operate on the same

```
#!/bin/bash
echo "Please enter your name "
read name
echo "hello $name , please enter city you are residing in "
read city
echo "$city is a good city "
```

- Above ways of getting input can be used to get away from hard coding the details
- Even though shell scripts are not as evolved as other programming languages considering these are meant for just scripting purposes , they do support looping as well
- As of today we will just see if else loop
- These are conditional loops which one can use execute simple condition based tasks

```
#!/bin/bash
clear
echo -e "What is your name : \c"
read name
echo hello $name. Welcome to Shell programming
sleep 2
clear
echo -e "Would you like to see a listing of the files in your Current Working Directory? [y/n]: \c"
read yn
if [ $yn = y ]
then
    ls
fi
```

```
#!/bin/bash
echo "enter the number "
read num
```

```

if [  $$(num\%2)$  -eq 0 ]
then
    echo "even number "
else
    echo "not even number "
fi

```

```

#!/bin/bash
read -p "Enter Current Year (YYYY): " input_year_val
current_year=`date +%Y`
if test $input_year_val == $current_year
then
    echo "You have entered the correct Year"
fi
if test $input_year_val != $current_year
then
    echo "You have not entered the Current Year. Please try again."
fi

```

```

#!/bin/bash
file=$1
echo "File path is $file"
if [ -f $file ] ; then
    echo "$file exists"
else
    echo "$file does not exist"
fi

```

./script.sh /etc/passwd

- as you can see above if loop begins with **if** and ends with **fi**
- things to be executed under then loop has to be properly intended
- same goes for else

```

echo "enter your access key "
read accesskey
echo "enter your secret access key "
read secret
echo "enter the bucket that you need backup in "
read bucket
echo "enter the backup folder full path "
read sourcepath
export AWS_ACCESS_KEY_ID=$accesskey
export AWS_SECRET_ACCESS_KEY=$secret
export AWS_DEFAULT_REGION=us-east-1
aws s3 cp $sourcepath s3://$bucket/backup/ --recursive

```

- above scripts puts the things we have learnt so far in perspective
- It will take path from user and then based on the inputs initiate the backup

```
#!/bin/bash
echo "enter your access key "
read accesskey
echo "enter your secret access key "
read secret
echo "enter the bucket that you need backup in "
read bucket
echo "enter the backup folder full path "
read sourcepath
export AWS_ACCESS_KEY_ID=$accesskey
export AWS_SECRET_ACCESS_KEY=$secret
export AWS_DEFAULT_REGION=us-east-1
aws s3 cp $sourcepath s3://$bucket/backup/ --recursive
if [ $? -eq 0 ]
then
    echo "backup successful"
else
    echo "backup failed , check your inputs "
fi
```

- To make the script better here , we are using `$?` operator . This stores the output of the previous command .
- If the output is 0 , previous command was successful , or else it failed
- Debugging a shell script
 - Debugging helps you troubleshoot and resolve such errors, and is one of the most important tasks a system administrator performs.

```
bash -x ./script_file
```

Scaling the storage capacity of the instance

- We know that we can change the instance type if we run out of computing capacity . In order to scale the storage capacity there are couple of approaches we can take
 - Scale the existing ebs volume
 - The general purpose ebs volume has upto 16tb of capacity . If a disc runs out of capacity , we can scale the existing volume

```
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
xvda        202:0    0  8G  0 disk
└─xvda1     202:1    0  8G  0 part /
[ec2-user@ip-172-31-16-9 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        482M   0  482M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M 456K  492M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.4G   6.7G  17% /
tmpfs           99M   0   99M   0% /run/user/1000
tmpfs           99M   0   99M   0% /run/user/0
[ec2-user@ip-172-31-16-9 ~]$
```

- As seen above , the root volume is of 8 gb . It is currently 17% used . We will now scale it
- Navigate to the aws console
- Go to the instance , and select the ebs volume attached to it . It navigates to the volumes screen . Click on actions and select modify volume
- Remember that we can only increase ebs volume .
- Add the new capacity and click on modify

Modify Volume ✕

Volume ID vol-029b2807e3d545239

Volume Type General Purpose SSD (gp2) ⓘ

Size (Min: 1 GiB, Max: 16384 GiB) ⓘ

IOPS 100 / 3000 (Baseline of 3 IOPS per GiB with a minimum of 100 IOPS, burstable to 3000 IOPS) ⓘ

Cancel Modify

- Once the volume is modified , come back to the terminal

```
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
xvda        202:0    0  8G  0 disk
└─xvda1     202:1    0  8G  0 part /
[ec2-user@ip-172-31-16-9 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        482M   0  482M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M 456K  492M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.4G   6.7G  17% /
tmpfs           99M   0   99M   0% /run/user/1000
tmpfs           99M   0   99M   0% /run/user/0
[ec2-user@ip-172-31-16-9 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        482M   0  482M   0% /dev
tmpfs           492M   0  492M   0% /dev/shm
tmpfs           492M 456K  492M   1% /run
tmpfs           492M   0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.4G   6.7G  17% /
tmpfs           99M   0   99M   0% /run/user/1000
tmpfs           99M   0   99M   0% /run/user/0
[ec2-user@ip-172-31-16-9 ~]$
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
xvda        202:0    0 20G  0 disk
└─xvda1     202:1    0  8G  0 part /
[ec2-user@ip-172-31-16-9 ~]$
```

- Notice that the volume is increased , but it wont be usable until we extend the filesystem

```
# based on the block id use
sudo growpart /dev/xvda 1
#above command will increase the partition to accomodate the new size
lsblk
#observe the changes from previous output
df -hT
# if you dont see any changes reflected in above command , use
sudo xfs_growfs -d /
df -hT
```

```
[ec2-user@ip-172-31-16-9 ~]$ sudo growpart /dev/xvda 1
CHANGED: partition=1 start=4096 old: size=16773087 end=16777183 new: size=41938911 end=41943007
```

```
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
xvda        202:0    0  20G  0 disk
└─xvda1    202:1    0  20G  0 part /
[ec2-user@ip-172-31-16-9 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        482M    0  482M   0% /dev
tmpfs           492M    0  492M   0% /dev/shm
tmpfs           492M  464K  492M   1% /run
tmpfs           492M    0  492M   0% /sys/fs/cgroup
/dev/xvda1      8.0G  1.4G  6.7G  17% /
tmpfs           99M    0   99M   0% /run/user/1000
tmpfs           99M    0   99M   0% /run/user/0
[ec2-user@ip-172-31-16-9 ~]$ sudo xfs_growfs -d /
meta-data=/dev/xvda1      isize=512    agcount=4, agsize=524159 blks
                =                  sectsz=512    attr=2, projid32bit=1
                =                  crc=1      finobt=1 spinodes=0
data        =                  bsize=4096   blocks=2096635, imaxpct=25
                =                  sunit=0      swidth=0 blks
naming      =version 2       bsize=4096   ascii-ci=0 ftype=1
log         =internal       bsize=4096   blocks=2560, version=2
                =                  sectsz=512   sunit=0 blks, lazy-count=1
realtime    =none           extsz=4096   blocks=0, rtextents=0
data blocks changed from 2096635 to 5242363
[ec2-user@ip-172-31-16-9 ~]$ df -hT
Filesystem      Type      Size  Used Avail Use% Mounted on
devtmpfs        devtmpfs  482M    0  482M   0% /dev
tmpfs           tmpfs     492M    0  492M   0% /dev/shm
tmpfs           tmpfs     492M  464K  492M   1% /run
tmpfs           tmpfs     492M    0  492M   0% /sys/fs/cgroup
/dev/xvda1      xfs       20G    1.4G   19G   7% /
tmpfs           tmpfs     99M    0   99M   0% /run/user/1000
tmpfs           tmpfs     99M    0   99M   0% /run/user/0
[ec2-user@ip-172-31-16-9 ~]$
```

- This way we can scale an existing volume
- Since the general purpose ebs has 16tb limit , some teams will run out of storage .
- In that scenario we can create new ebs volume and attach it to the instance
- In order to do that , first we need to create a volume
- Navigate to the volumes screen in ec2 service
- Click on create new volume
- Make sure this volume is created in the exact same AZ as the instance
- Once this volume is created attach it the instance
- Notice the changes by using "lsblk" command

```
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
xvda        202:0    0   20G  0 disk
└─xvda1     202:1    0   20G  0 part /
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
xvda        202:0    0   20G  0 disk
└─xvda1     202:1    0   20G  0 part /
xvdf        202:80   0  100G  0 disk
[ec2-user@ip-172-31-16-9 ~]$
```

- Notice that the mountpoint section in front of newly added block device is blank . Which means it is not usable yet
- First we will have to create a file system on the volume . Remember , without file system we cannot store data on any disc

```
# first let us check if the volume has a file system
sudo file -s /dev/xvdf
# If output of above command comes as "data", then it means it does not have a
file system and we can go ahead and create it .
# However if the output is x86 or xfs , it means it already has a filesystem and
we shouldnt create it . If you create a filesystem on this it will erase all the
existing data . SO make sure to check before you proceed
sudo mkfs -t xfs /dev/xvdf
# Above command will create a file system on the disc
#now we need a directory where the disc needs to be mounted. It is basically the
address where the disc will be utilized
sudo mkdir /data
sudo mount /dev/xvdf /data
```

```
[ec2-user@ip-172-31-16-9 ~]$ sudo file -s /dev/xvdf
/dev/xvdf: data
```

```
[ec2-user@ip-172-31-16-9 ~]$ sudo mkfs -t x86 /dev/xvdf
mkfs: failed to execute mkfs.x86: No such file or directory
[ec2-user@ip-172-31-16-9 ~]$ sudo mkfs -t xfs /dev/xvdf
meta-data=/dev/xvdf            isize=512    agcount=4, agsize=6553600 blks
=                               sectsz=512   attr=2, projid32bit=1
=                               crc=1        finobt=1, sparse=0
data      =                     bsize=4096  blocks=26214400, imaxpct=25
=                               sunit=0        swidth=0 blks
naming    =version 2           bsize=4096  ascii-ci=0 ftype=1
log       =internal log       bsize=4096  blocks=12800, version=2
=                               sectsz=512   sunit=0 blks, lazy-count=1
realtime  =none               extsz=4096  blocks=0, rtextents=0
[ec2-user@ip-172-31-16-9 ~]$
```

```
[ec2-user@ip-172-31-16-9 ~]$ sudo mkdir /data
[ec2-user@ip-172-31-16-9 ~]$ sudo mount /dev/xvdf /data
[ec2-user@ip-172-31-16-9 ~]$ lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINT
xvda        202:0    0   20G  0 disk
└─xvda1     202:1    0   20G  0 part /
xvdf        202:80   0  100G  0 disk /data
[ec2-user@ip-172-31-16-9 ~]$
```

- Now the directory and the new volume are usable
- This mounted directory however will be reset if the instance goes through a reboot
- In order to avoid that , we need to enter these new changes in fstab file
- The entry requires a UUID , which is globally unique


```
sudo blkid
```

above command will give you UUID of all the block devices attached to the machine. Copy for the newly created volume

Now we will make the changes the fstab file. fstab file is one of the most crucial files in linux, if there are any mistakes in the entry , then after the next reboot , you will not be able to login on the instance . So make the changes carefully , you can also take backup of the current config by creating a copy of it

```
sudo vi /etc/fstab
```

#save and exit

in order to confirm if all the entries in fstab are correct try below command

```
sudo mount -a
```

If above command runs successfully , then the entry syntax is correct . You can try rebooting the instance and verify the same

```
[ec2-user@ip-172-31-16-9 ~]$ sudo cat /etc/fstab
#
UUID=15c7809d-e6e3-4062-a5eb-afeb1939fc6e    /          xfs     defaults,noatime 1 1
[ec2-user@ip-172-31-16-9 ~]$ sudo blkid
/dev/xvda1: LABEL="/" UUID="15c7809d-e6e3-4062-a5eb-afeb1939fc6e" TYPE="xfs" PARTLABEL="Linux" PARTUUID="cfb2e895-ecec-41c5-afc2-40ffe2e4384a"
/dev/xvdf: UUID="434aa9a1-f68c-4254-af5a-1bc8df9f35f3" TYPE="xfs"
[ec2-user@ip-172-31-16-9 ~]$ sudo vi /etc/fstab
[ec2-user@ip-172-31-16-9 ~]$
```

```
#
UUID=15c7809d-e6e3-4062-a5eb-afeb1939fc6e    /          xfs     defaults,noatime 1 1
UUID=434aa9a1-f68c-4254-af5a-1bc8df9f35f3    /data      xfs     defaults,noatime 1 1
~
~
~
~
~
~
~
~
```